

Near-term Ecological Forecasting Initiative Summer Course

Boston University
Boston, MA 02215

June 8 - June 12, 2020

<https://ecoforecast.org/nefi2020/>

Course Description

Iterative ecological forecasts are continually-updated, actionable predictions, with uncertainties, of the state of ecosystems and their services. The backbone of this course covers topics related to the statistics of model-data fusion and forecasting: Bayesian statistics; machine learning; uncertainty partitioning, propagation, and analysis; fusing multiple data sources; assessing model performance; and a suite of iterative data assimilation techniques. While the goals of some ecological forecasts are simply to advance basic science, many are performed to inform the public, managers, and decision makers. Therefore the course also covers the basics of structured decision making and expert elicitation.

Course Objectives

After taking this course, you should be able to generate a forecast from a simple model that makes use of more than one data stream. You should be able to categorize and estimate the uncertainties in the model, propagate this uncertainty into the forecast, and assess the performance of the model. You should be able to assess and critique forecast data products and scientific publications about ecological forecasts, and understand how such forecasts can be used to inform decision making.

Textbook

Dietze, M. C. (2017). Ecological Forecasting. Princeton, NJ: Princeton University Press.

Course Organizers

Name	Institution	email	cell
Mike Dietze, PI	Boston Univ.	dietze@bu.edu	410-533-5130
Shannon LaDeau, coPI	Cary Institute	ladeaus@caryinstitute.org	919-323-5983
Jenny Bhatnagar, coPI	Boston Univ.	jmtalbot@bu.edu	
Kathie Weathers, coPI	Cary Institute	weathersk@caryinstitute.org	
Quinn Thomas	Virginia Tech	rqthomas@vt.edu	
Kira Sullivan-Wiley	Boston Univ	kira.s.wiley@gmail.com	
Ethan Deyle	Boston Univ	edeyle@bu.edu	
Jason McLachlan	Notre Dame	Jason.S.McLachlan.2@nd.edu	
Jody Peters	Notre Dame	peters.63@nd.edu	
John Foster, TA	Boston Univ.	fosterj@bu.edu	651-558-7362
Kathryn Wheeler, TA	Boston Univ.	kiwheel@bu.edu	828-773-9737
Zoey Werbin, TA	Boston Univ.	zrwerbin@bu.edu	954-496-2100
Katie Zarada, TA	Boston Univ	kzarada@bu.edu	912-660-1871

Thursday Round Table

Discussion of issues that span academia, agencies, and non-traditional research both generally (e.g. career options, publishing, funding, etc) and in the specific context of trying to build a community of practice in ecological forecasting. Where is ecological forecasting going, how do we get it there?

Schedule

Time	Monday	Tuesday	Wednesday	Thursday	Friday
9:00	Introductions (0:40)	State-space (0:30)	Propagating Uncertainty (1:15)	Ensemble DA (1:15)	Model Assessment (1:00)
9:30	Lightning talks (0:35)	Hands On			
10:00	Project talks (0:20)		Break (0:15)	Break (0:15)	Break (0:30)
10:30	Break	Break	Hands On (1:30)	Hands On (1:15)	Project
11:00	Characterizing Uncertainty (30 min)	Dynamic Models (0:30)			
11:30	Hands On	Project		Break (0:15)	
12:00			Lunch	Social Sci (0:30)	
12:30	Lunch	Lunch		Lunch	Lunch
13:00			Analytical DA (1:15)		
13:30	Hierarchical Bayes (1:15)			Project	Forecast Infrastructure / FLARE case study (1:15)
14:00			Break (0:15)		
14:30	Break (0:15)	Break (0:15)	Hands On (1:15)		Break (0:15)
15:00	hands on (1:15)	Machine Learning (1:00)			Project Presentations (1:00)
15:30		Break (0:15)	Break (0:15)	Break (0:15)	
16:00	Break (0:15)	PROACT (1:00)	Project	Round Table (1:00)	Wrap up (0:45)
16:30	Expert Elicitation (0:30)				

Hands-on Activities

The goal of the course isn't just to expose you to the concepts of ecological forecasting, but to give you hands-on experience in applying the tools used. Therefore most lectures are paired with hands-on activities. Most of these activities are independent (you don't need to complete one to move on to the next) and many are open-ended -- there's more you could play with than we expect you to be able to complete in the allocated time. The instructors will often state what the minimal goal of each activity is. Beyond that, we encourage you to complete what you can, play with the activities, and know that you can come back to the activities that you find most important and interesting at a later date. We will split you into randomized break out groups for each of the hands on activities. Each group will have TAs and instructors to assist with the activity.

Group Projects

The goal of the group project is to take what you've learned in lecture and the hands-on activities and see if you can apply it to a new problem. Students will work in small groups for the final project and make a 10-15 min presentation about what they've accomplished this week on Friday afternoon. The goals of the project are:

- **Calibration:** Fit a simple dynamic model to time-series data, preferably using a Bayesian state-space framework and accounting for the complexities of the data. Make sure to hold out some portion of the available data for validation.
- **Forecast:** Predict or project the dynamics of the system forward from the calibration period into the validation period while accounting for uncertainties. If time permits, partition the forecast uncertainties to determine which dominates.
- **Validation:** Assess model performance against held-out data
- **Analysis:** For the next time point beyond the calibration period, use data assimilation to update the state of the system and re-forecast. If time permits, iteratively assimilate all validation data. How does the uncertainty in your forecast change with lead time and what is the forecast horizon for this system?

The set of projects available to participate in has been pre-selected based on the participation of federal agency personnel as part of the course. The goal is to apply forecasting principles to real-world agency problems, to help agency participants get a jump start on implementing forecasts, and to help establish foster forecasting collaborations between academia and agencies. The set of available projects will be described Monday morning. Everyone will be asked to rank their interest in the available projects and then will be assigned to a team by the instructors

R / RStudio

The course activities will be done in R and will be provided as R Markdown (Rmd) documents. Use of RStudio is not required, but activities are prepared assuming you will be working in that IDE.

Git / Github

Course activities and in-house R packages will be distributed via Github, so at a minimum you will be expected to be able to clone / download those repositories (we'll be sending a follow up email in a few weeks with more details on exactly what repositories to install). Bare-bones basics on how to use Github to retrieve course materials is located at the top of our basic R tutorial: https://github.com/EcoForecast/EF_Activities/blob/master/Exercise_01_RPrimer.Rmd

While not required just to download, you'll probably want to create an account at <https://github.com/> if you don't have one already.

You are also strongly encouraged to use Github (or equivalent) when collaborating on code for your course projects. We also have a tutorial on using Github for collaborative coding, which is

best completed with a partner:

https://github.com/EcoForecast/EF_Activities/blob/master/Exercise_04_PairCoding.Rmd

Code & data management

These are covered in chapters 3 and 4 of the book, which we won't be covering in class.

You may also enjoy:

Wilson, G., Bryan, J., Cranston, K., Kitzes, J., Nederbragt, L., & Teal, T. K. (2016). Good Enough Practices in Scientific Computing. PLoS Computational Biology, 13, 1–20.

<http://doi.org/10.1371/journal.pcbi.1005510>

Bayesian statistics

In addition to basic concepts, we're also going to assume basic familiarity with Bayesian software. Specifically, we will be using JAGS (called through R) for the course tutorials, so you will want to make sure you have install the JAGS software ahead of time (<http://mcmc-jags.sourceforge.net/>). As a point of reference, we'd like everyone to be able to implement a basic Bayesian linear regression model in JAGS (this is the milestone we aim to reach in the optional Bayes primer).

Software

As mentioned previously we'll be using R / RStudio, JAGS, and Git / Github so you'll want to make sure that those are installed on your laptop prior to the course.

Furthermore, all the course activities are in the following repository:

https://github.com/EcoForecast/EF_Activities

Which you'll want to 'git clone' (or fork to your local repo) before the course starts. A number of activities do rely on additional packages. Most of these are quick & easy, but do make sure you have 'rjags' installed and working.

Readings

From the *Ecological Forecasting* textbook, we're going to assume Ch 3-5 as prior knowledge and will focus the course on Chapters 2, 6, 8, 9, 11, 13, 14, and 17. We're not requiring that folks have read these chapters ahead of time, but if you do want to start diving in ahead of time that's where we recommend you focus. Any advanced reading you're able to do will make following lectures and doing the activities that much easier.

In addition to the textbook we've provided the following readings as additional references:

Munch, Giron-Nava, and Sugihara, 2018. "Nonlinear Dynamics and Noise in Fisheries Recruitment." Fish and Fisheries 19(6): 964-973 <https://doi.org/10.1111/faf.12304>

Ye H, RJ Beamish, SM Glaser, SCH Grant, C-H Hsieh, LJ Richards, JT Schnute, G Sugihara. 2015. Equation-free ecosystem forecasting. *Proceedings of the National Academy of Sciences* 112 (13) E1569-E1576; DOI: 10.1073/pnas.1417063112

Morgan, M. Granger. 2014. Expert elicitation for decision making. *Proceedings of the National Academy of Sciences*, 111 (20) 7176-7184; DOI: 10.1073/pnas.1319946111